

# lab-1

October 26, 2024

## 0.0.1 1.1 Variable Classification

Based on the dataset, we classify each variable into one of the following categories: 1. **Nominal Categorical**: Categories with no specific order. 2. **Binary Categorical**: Two possible categories, such as yes/no or 0/1. 3. **Discrete**: Whole numbers, like a count of items. 4. **Continuous**: Values that can take any value within a range, such as height or temperature.

Let's load the data and assign each column to a category.

```
[28]: # Importing necessary libraries
import pandas as pd

# Load the dataset
df = pd.read_csv('heart.csv')

# Categorize each variable
variable_classification = {
    'sbp': 'Continuous',
    'tobacco': 'Continuous',
    'ldl': 'Continuous',
    'adiposity': 'Continuous',
    'famhist': 'Binary Categorical',
    'typea': 'Continuous',
    'obesity': 'Continuous',
    'alcohol': 'Continuous',
    'age': 'Discrete',
    'chd': 'Binary Categorical'
}

# Display the classification
for column, category in variable_classification.items():
    print(f"{column}: {category}")
```

```
sbp: Continuous
tobacco: Continuous
ldl: Continuous
adiposity: Continuous
famhist: Binary Categorical
typea: Continuous
obesity: Continuous
```

alcohol: Continuous  
age: Discrete  
chd: Binary Categorical

### 0.0.2 1.2 Find the Number of Null Values for Each Column

In this section, we will check for any missing values in each column of the dataset and display the count of null values.

```
[31]: # Check for null values in each column
null_counts = df.isnull().sum()

# Display the number of null values for each column
print("Number of null values for each column:")
print(null_counts)
```

Number of null values for each column:

```
sbp      28
tobacco  40
ldl      39
adiposity 40
famhist  45
typea    41
obesity  40
alcohol  40
age       35
chd       39
dtype: int64
```

### 0.0.3 1.31 General Descriptive Statistics

Let's display the general descriptive statistics of the dataset to understand the basic distribution of each variable.

```
[34]: # Display general descriptive statistics
print("General Descriptive Statistics:")
df.describe()
```

General Descriptive Statistics:

```
[34]:
```

	sbp	tobacco	ldl	adiposity	typea	obesity	\
count	384.000000	372.000000	373.000000	372.000000	371.000000	372.000000	
mean	139.216146	3.676425	4.569303	25.210753	52.008086	25.763602	
std	20.307368	4.568564	1.888691	7.760257	9.822888	3.854265	
min	101.000000	0.000000	0.980000	7.120000	20.000000	17.890000	
25%	124.000000	0.057500	3.240000	19.307500	46.000000	22.835000	
50%	136.000000	1.800000	4.220000	26.115000	52.000000	25.675000	
75%	148.500000	5.640000	5.470000	30.790000	58.000000	28.167500	
max	218.000000	27.400000	14.160000	42.490000	73.000000	40.340000	

	alcohol	age	chd
count	372.000000	377.000000	373.000000
mean	18.425134	42.453581	0.335121
std	25.971090	15.312649	0.472667
min	0.000000	15.000000	0.000000
25%	0.195000	30.000000	0.000000
50%	7.300000	45.000000	0.000000
75%	25.820000	57.000000	1.000000
max	145.290000	64.000000	1.000000

#### 0.0.4 1.32 Oldest Person

Identify the age of the oldest person in the dataset and display all individuals who are of that age.

```
[37]: # Find the maximum age and filter the dataset
oldest_age = df['age'].max()
oldest_people = df[df['age'] == oldest_age]

print(f"The age of the oldest person is: {oldest_age}")
print("People with the oldest age:")
oldest_people
```

The age of the oldest person is: 64.0

People with the oldest age:

```
[37]:
```

	sbp	tobacco	ldl	adiposity	famhist	typea	obesity	alcohol	age \
58	158.0	3.60	2.97	NaN	Absent	NaN	26.64	108.00	64.0
70	152.0	12.18	4.04	37.83	Present	63.0	34.57	4.17	64.0
110	126.0	0.00	5.98	29.06	Present	56.0	25.39	11.52	64.0
167	148.0	8.20	7.75	34.46	Present	46.0	26.53	6.04	64.0
170	128.0	5.16	4.90	NaN	Present	57.0	26.42	0.00	64.0
206	NaN	8.60	3.90	32.16	Present	52.0	28.51	11.11	64.0
241	160.0	0.60	6.94	30.53	Absent	36.0	25.68	1.42	64.0
256	138.0	2.00	5.11	31.40	Present	49.0	27.25	2.06	64.0
276	128.0	0.73	3.97	23.52	Absent	NaN	23.81	NaN	64.0
348	140.0	8.60	3.90	32.16	Present	52.0	28.51	11.11	64.0
374	160.0	0.60	6.94	30.53	Absent	36.0	25.68	NaN	64.0
402	174.0	2.02	6.57	31.90	Present	50.0	28.75	11.83	64.0

	chd
58	0.0
70	0.0
110	1.0
167	1.0
170	0.0
206	1.0
241	0.0

```
256  1.0
276  0.0
348  1.0
374  0.0
402  1.0
```

### 0.0.5 1.33 Youngest Person

Identify the age of the youngest person in the dataset and display all individuals who are of that age.

```
[40]: # Find the minimum age and filter the dataset
youngest_age = df['age'].min()
youngest_people = df[df['age'] == youngest_age]

print(f"The age of the youngest person is: {youngest_age}")
print("People with the youngest age:")
youngest_people
```

The age of the youngest person is: 15.0

People with the youngest age:

```
[40]:      sbp  tobacco   ldl  adiposity famhist  typea  obesity  alcohol  age \
9    132.0      0.0  1.87    17.21  Absent   49.0    23.63    0.97  15.0
38     NaN      0.0  3.67    12.13  Absent    NaN    19.15    0.60  15.0

      chd
9    0.0
38   0.0
```

### 0.0.6 1.34 Average and Standard Deviation of Age

Calculate the average (mean) and standard deviation of the age column to understand its central tendency and spread.

```
[43]: # Calculate mean and standard deviation of the age column
age_mean = df['age'].mean()
age_std = df['age'].std()

print(f"Average age: {age_mean}")
print(f"Standard deviation of age: {age_std}")
```

Average age: 42.45358090185676

Standard deviation of age: 15.31264927550187

### 0.0.7 1.35 Median Age

Calculate the median of the age column.

```
[46]: # Calculate the median age
age_median = df['age'].median()

print(f"Median age: {age_median}")
```

Median age: 45.0

### 0.0.8 1.36 Bar Chart of Deaths vs. Age


Filter the data for individuals who have died ( $chd = 1$ ) and create a histogram to observe the age distribution of deaths.

```
[49]: import matplotlib.pyplot as plt

# Filter data for deaths (chd = 1)
deaths_data = df[df['chd'] == 1]

# Plot histogram
plt.figure(figsize=(10, 6))
plt.hist(deaths_data['age'], bins=10, color='skyblue', edgecolor='black')
plt.xlabel("Age")
plt.ylabel("Number of Deaths")
plt.title("Age Distribution of Deaths (chd = 1)")
plt.show()

# Insight
# From this chart, we can observe that ...
```



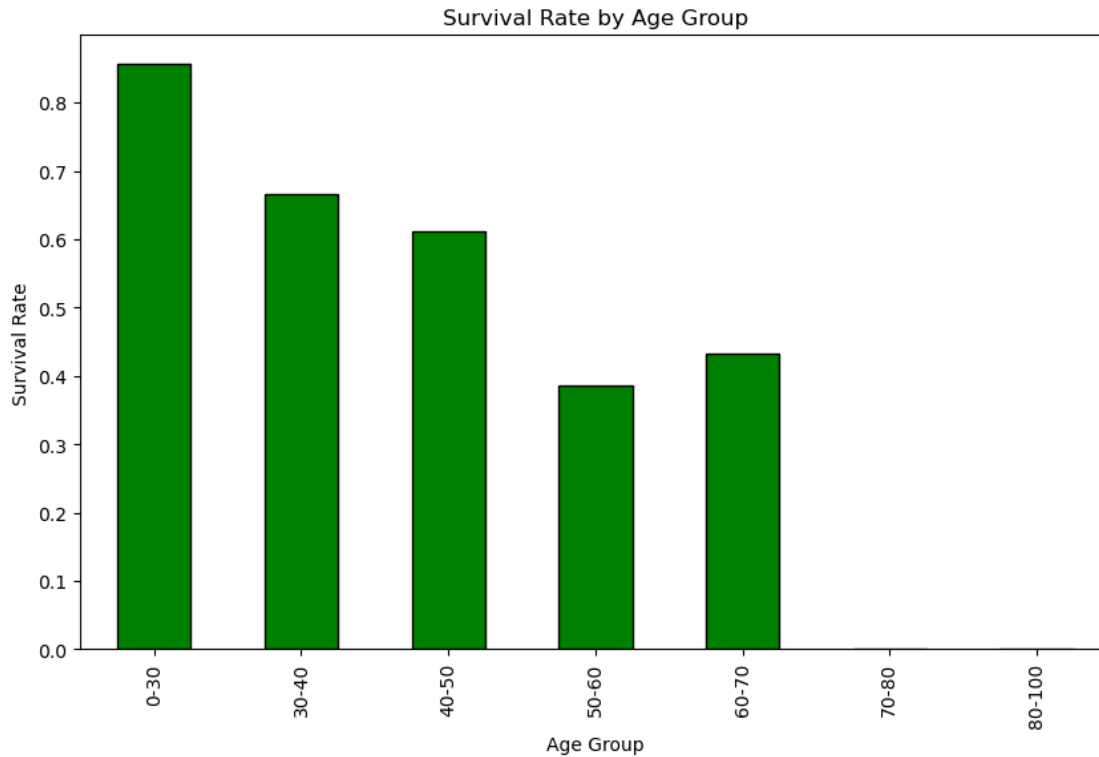
### 0.0.9 1.37 Age Groups with Largest Survival Rate

Divide the data into age groups and calculate the survival rate for each group, then visualize it.

```
[52]: # Define age groups and calculate survival rates
df['age_group'] = pd.cut(df['age'], bins=[0, 30, 40, 50, 60, 70, 80, 100],
    ↳ labels=['0-30', '30-40', '40-50', '50-60', '60-70', '70-80', '80-100'])
survival_rate = df[df['chd'] == 0].groupby('age_group').size() / df.
    ↳ groupby('age_group').size()

# Plot survival rate
plt.figure(figsize=(10, 6))
survival_rate.plot(kind='bar', color='green', edgecolor='black')
plt.xlabel("Age Group")
plt.ylabel("Survival Rate")
plt.title("Survival Rate by Age Group")
plt.show()
```

```
C:\Users\Niranj Rao\AppData\Local\Temp\ipykernel_13800\1887164646.py:3:
FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    survival_rate = df[df['chd'] == 0].groupby('age_group').size() /
df.groupby('age_group').size()
C:\Users\Niranj Rao\AppData\Local\Temp\ipykernel_13800\1887164646.py:3:
FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    survival_rate = df[df['chd'] == 0].groupby('age_group').size() /
df.groupby('age_group').size()
```



#### 0.0.10 1.38 Relationship between Family History (famhist) and Coronary Heart Disease (chd)

Analyze the relationship between famhist and chd.

```
[55]: # Calculate the relationship between famhist and chd
famhist_chd_relationship = df.groupby('famhist')['chd'].mean()

print("Relationship between Family History and CHD:")
print(famhist_chd_relationship)
```

Relationship between Family History and CHD:

famhist

Absent 0.243094

Present 0.443709

Name: chd, dtype: float64

#### 0.0.11 1.39 Visualizations for Data Distributions

Let's create various visualizations to better understand data distributions.

```
[70]: import seaborn as sns
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
```

```

# i. Correlation Matrix
plt.figure(figsize=(10, 8))
numeric_df = df.select_dtypes(include=['float64', 'int64']) # Select only
↳ numeric columns
sns.heatmap(numeric_df.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Matrix")
plt.show()


# ii. Scatter Matrix
plt.figure(figsize=(12, 12))
scatter_matrix(numeric_df, figsize=(12, 12), diagonal='kde')
plt.suptitle("Scatter Matrix")
plt.show()

# iii. Per Column Distribution
plt.figure(figsize=(12, 10))
numeric_df.hist(bins=20, edgecolor='black', figsize=(12, 10))
plt.suptitle("Per Column Distribution")
plt.tight_layout()
plt.show()

# iv. Heatmap for Missing Values
plt.figure(figsize=(12, 6))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title("Heatmap of Missing Values")
plt.show()

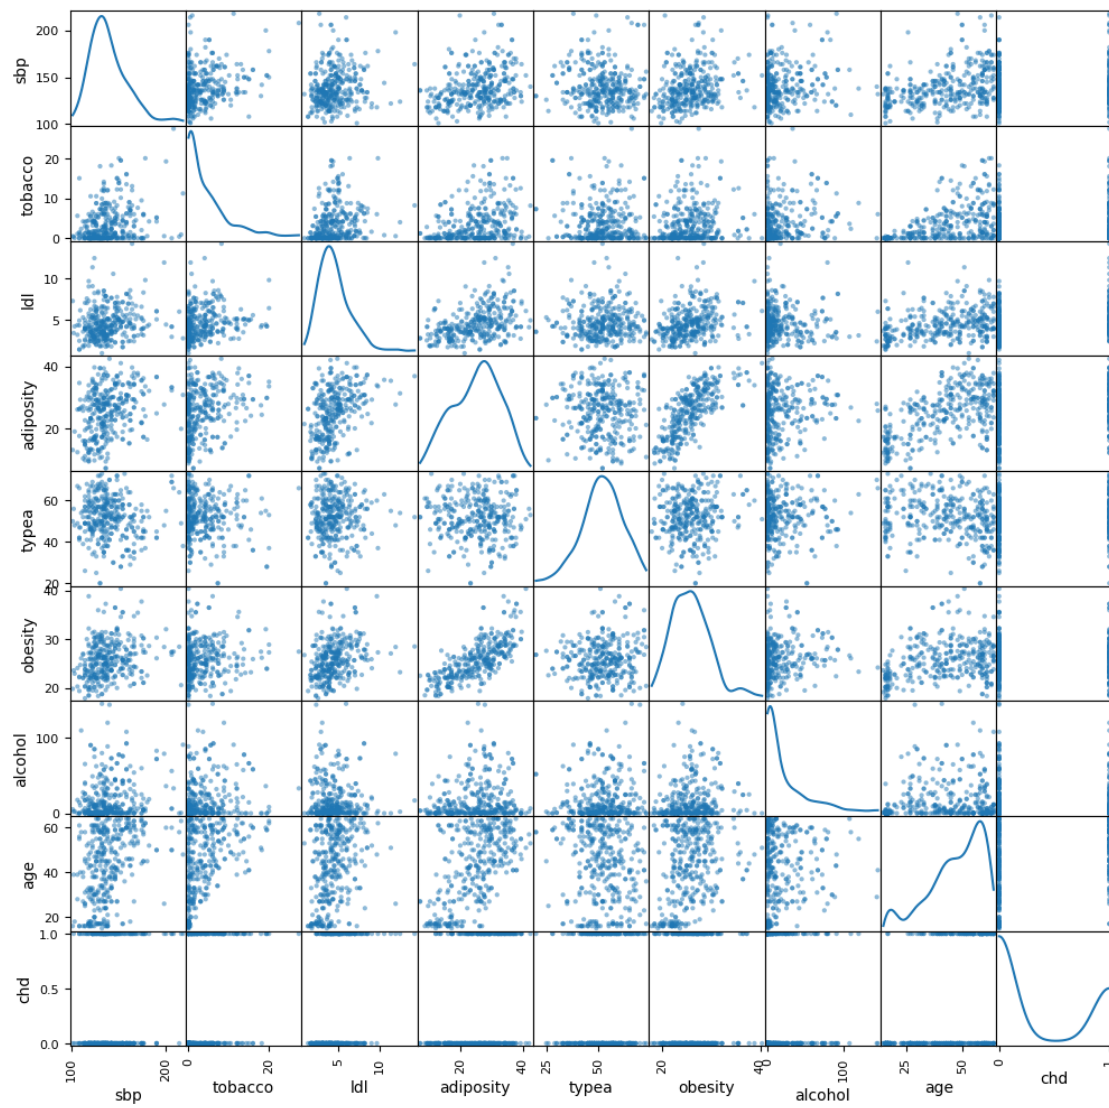
```




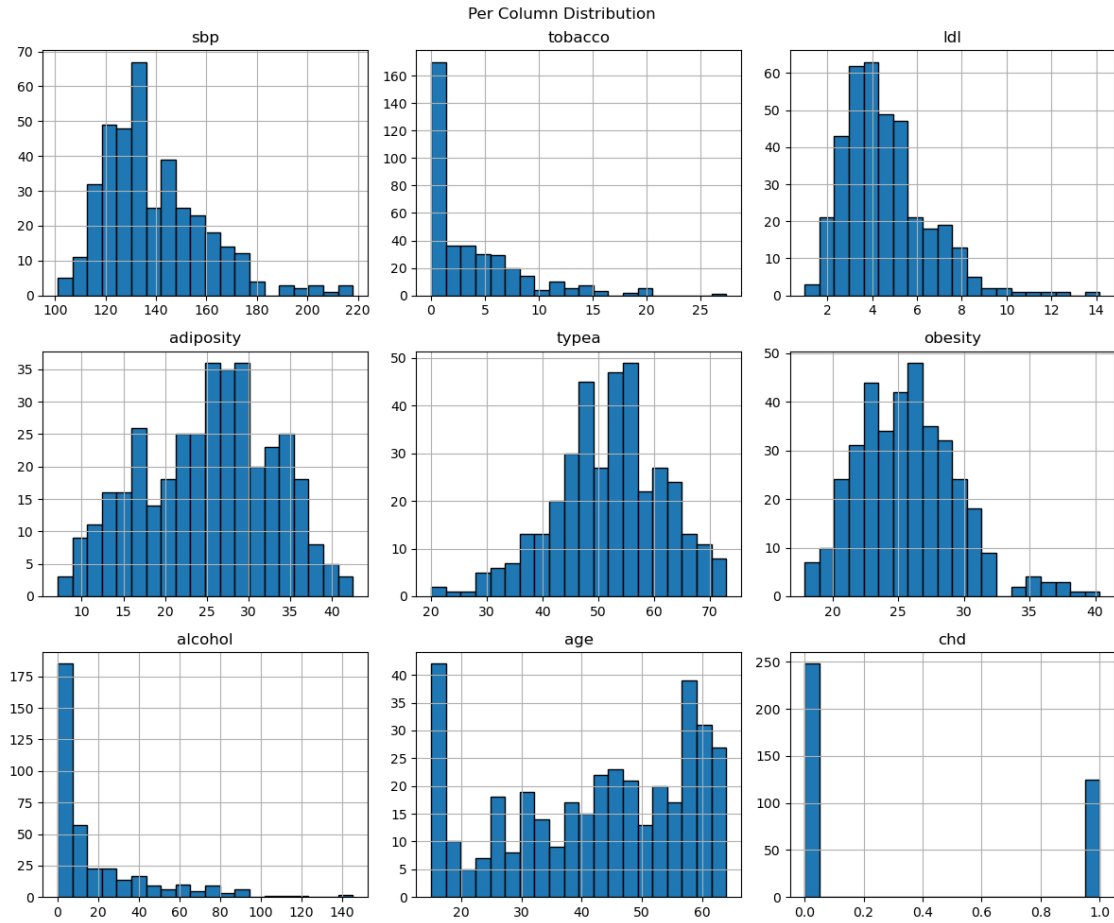


<Figure size 1200x1200 with 0 Axes>

Scatter Matrix



<Figure size 1200x1000 with 0 Axes>



### 0.0.12 1.3.10 Handling Null Values

Aside from simply dropping rows with missing data, there are several techniques for handling null values:

1. **Imputation:**
  - Replace missing values with the mean, median, or mode of the column.
  - Useful for continuous variables (e.g., replacing with mean or median) or categorical variables (e.g., replacing with mode).
2. **Interpolation:**
  - For time series data, interpolate missing values based on neighboring values, maintaining data continuity.
3. **Using Machine Learning Models:**
  - Train a predictive model to estimate missing values based on other features in the dataset.
  - Examples include k-Nearest Neighbors (k-NN) and linear regression.
4. **Using Domain-Specific Values:**
  - For some fields, domain knowledge can inform a reasonable substitute for missing values.
5. **Multiple Imputation:**
  - Create multiple imputations for missing values and use the resulting datasets for robust statistical analysis.

Choosing the best approach depends on the type of data, its distribution, and the analysis goals.

### 0.0.13 2.1 Basic Matrix Multiplication

Define a function that multiplies two matrices, A and B, without using any built-in matrix multiplication functions.

```
[85]: import numpy as np

def matrix_multiply(A, B):
    """
    Multiplies two matrices A and B without using built-in matrix multiplication
    functions.

    Parameters:
    A (numpy.ndarray): First matrix.
    B (numpy.ndarray): Second matrix.

    Returns:
    numpy.ndarray: The product of matrices A and B.
    """
    # Check if the matrices can be multiplied
    if A.shape[1] != B.shape[0]:
        raise ValueError("Number of columns in A must equal the number of rows
        in B.")
```

```

# Initialize the result matrix with zeros
result = np.zeros((A.shape[0], B.shape[1]))

# Perform matrix multiplication
for i in range(A.shape[0]):
    for j in range(B.shape[1]):
        for k in range(A.shape[1]):
            result[i][j] += A[i][k] * B[k][j]

return result

# Example usage
A = np.array([[12, 21], [2, 8]])
B = np.array([[13, 7], [7, 8]])
print("Matrix A:")
print(A)
print("Matrix B:")
print(B)
print("Product of A and B:")
print(matrix_multiply(A, B))

```

```

Matrix A:
[[12 21]
 [ 2  8]]
Matrix B:
[[13  7]
 [ 7  8]]
Product of A and B:
[[303. 252.]
 [ 82.  78.]]

```

## 0.0.14 2.2 Compute the Determinant

Define a function to compute the determinant of a square matrix using the `numpy.linalg` module.

```

[88]: import numpy as np

def compute_determinant(A):
    """
    Computes the determinant of a square matrix A.

    Parameters:
    A (numpy.ndarray): Square matrix.

    Returns:
    float: Determinant of the matrix A.
    """
    return np.linalg.det(A)

```

```

# Example usage
A = np.array([[11, 13], [15, 17]])
print("Matrix A:")
print(A)
print("Determinant of A:")
print(compute_determinant(A))

```

```

Matrix A:
[[11 13]
 [15 17]]
Determinant of A:
-8.0000000000000012

```

### 0.0.15 2.3 Solve a System of Linear Equations

Define a function to solve the system of linear equations ( $Ax = b$ ) using `numpy.linalg.solve()`.

```

[68]: import numpy as np

def solve_linear_system(A, b):
    """
    Solves the system of linear equations  $Ax = b$ .

    Parameters:
    A (numpy.ndarray): Coefficient matrix.
    b (numpy.ndarray): Constant vector.

    Returns:
    numpy.ndarray: Solution vector x.
    """
    return np.linalg.solve(A, b)

# Example usage
A = np.array([[3, 1], [1, 2]])
b = np.array([9, 8])
print("Coefficient matrix A:")
print(A)
print("Constant vector b:")
print(b)
print("Solution vector x:")
print(solve_linear_system(A, b))

```

```

Coefficient matrix A:
[[3 1]
 [1 2]]
Constant vector b:
[9 8]
Solution vector x:

```

[2. 3.]