

Niranjan-Rao-Lab2

November 6, 2024

1 Q1: Recommender System (20 Points)

```
[100]: import pandas as pd
import numpy as np
from scipy.sparse.linalg import svds
from sklearn.metrics.pairwise import cosine_similarity
```

2 1. Load the movies and ratings data. (2 points)

```
[101]: movies = pd.read_csv('movies.dat', sep='::', header=None, names=['MovieID', 'Title', 'Genres'], engine='python', encoding='ISO-8859-1')
ratings = pd.read_csv('ratings.dat', sep='::', header=None, names=['UserID', 'MovieID', 'Rating', 'Timestamp'], engine='python', encoding='ISO-8859-1')
```

```
[102]: # Display the first few rows of the data for verification
print("Movies Data:")
print(movies.head())
print("\nRatings Data:")
print(ratings.head())
```

Movies Data:

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

Ratings Data:

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

3 2. What is Singular Value Decomposition (SVD)? Explain it in your own words. (2 points)

- Singular Value Decomposition (SVD) is a matrix factorization technique used to decompose a matrix into three component matrices: U , S , and V^T .
- U contains the left singular vectors, S is a diagonal matrix with singular values, and V^T contains the right singular vectors.
- This method is often used in recommendation systems to reduce dimensionality and find patterns in data.

4 3. Explain content-based vs collaborative recommendation. (2 points)

- Content-based recommendation systems suggest items similar to those the user has interacted with, based on item features.
- Collaborative recommendation systems use past interactions from multiple users to make suggestions, leveraging similarities between users or items.

5 4. Create $m \times u$ matrix with movies as rows and users as columns. Normalize the matrix. (2 points)

```
[103]: # Create user-item matrix
user_movie_matrix = ratings.pivot(index='MovieID', columns='UserID',
    ↪values='Rating').fillna(0)

# Normalize the matrix by subtracting the mean of each row
movie_mean = user_movie_matrix.mean(axis=1)
matrix_normalized = user_movie_matrix.subtract(movie_mean, axis=0)

print("User-Movie Matrix (First 5 rows):")
print(user_movie_matrix.head())
```

User-Movie Matrix (First 5 rows):

UserID	1	2	3	4	5	6	7	8	9	10	...	\
MovieID											...	
1	5.0	0.0	0.0	0.0	0.0	4.0	0.0	4.0	5.0	5.0	...	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	...	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	3.0	0.0	0.0	...	
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	

UserID	6031	6032	6033	6034	6035	6036	6037	6038	6039	6040
MovieID										
1	0.0	4.0	0.0	0.0	4.0	0.0	0.0	0.0	0.0	3.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0

4	0.0	0.0	0.0	0.0	2.0	2.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 6040 columns]

6 5. Perform SVD to get U, S, and V. (4 points)

```
[104]: from scipy.sparse import csr_matrix

# Convert the normalized matrix to a sparse format
matrix_normalized_sparse = csr_matrix(matrix_normalized.values)

# Perform SVD on the sparse matrix
U, S, Vt = svds(matrix_normalized_sparse, k=50) # k represents the number of
↪ components
S = np.diag(S)

print("Shapes after SVD:")
print(f"U: {U.shape}, S: {S.shape}, Vt: {Vt.shape}")
```

Shapes after SVD:

U: (3706, 50), S: (50, 50), Vt: (50, 6040)

7 6. Select top 30 components from S. (2 points)

```
[105]: # Select the top 30 singular values
U_30 = U[:, :30]
S_30 = S[:30, :30]
Vt_30 = Vt[:30, :]

print("Selected top 30 components from S.")
```

Selected top 30 components from S.

8 7. Get the top 30 eigenvectors using eigenvalues. (2 points)

- Eigenvectors corresponding to the largest eigenvalues are selected.

```
[106]: eigenvectors = U_30
print("Top 30 eigenvectors obtained.")
```

Top 30 eigenvectors obtained.

9 8. Using cosine similarity, find 10 closest movies using the 30 components from SVD. (2 points)

- Compute cosine similarity between movies

```
[107]: # Compute cosine similarity between movies
movie_similarity = cosine_similarity(U_30)

# Function to get the 10 most similar movies
def find_similar_movies(movie_id, num_similar=10):
    movie_idx = movies.index[movies['MovieID'] == movie_id].tolist()[0]
    similarities = list(enumerate(movie_similarity[movie_idx]))
    sorted_similarities = sorted(similarities, key=lambda x: x[1], reverse=True)
    similar_movies = [movies.iloc[i][0]]['Title'] for i in sorted_similarities[1:
    ↪ num_similar + 1]]
    return similar_movies

# Example: Find 10 closest movies to a specific movie
example_movie_id = movies['MovieID'].iloc[0] # Replace with any MovieID for ↪
    ↪ actual search
closest_movies = find_similar_movies(example_movie_id)
print(f"10 closest movies to '{movies[movies['MovieID'] == ↪
    ↪ example_movie_id]['Title'].values[0]}' are:")
print(closest_movies)
```

10 closest movies to 'Toy Story (1995)' are:
['Bad Seed, The (1956)', 'Barb Wire (1996)', 'First Wives Club, The (1996)',
'Entrapment (1999)', 'Babe (1995)', 'Action Jackson (1988)', 'Belizaire the
Cajun (1986)', 'Night Tide (1961)', 'Crow, The (1994)', 'Like Water for
Chocolate (Como agua para chocolate) (1992)']

10 9. Discuss results of above SVD methods. (2 points)

- The SVD method helps reduce dimensionality, allowing us to find hidden relationships between movies.
- The top 30 eigenvectors capture the most significant patterns, and cosine similarity helps find similar movies.
- This approach balances accuracy and computation efficiency, showing how a smaller subset of features can represent the dataset effectively.

11 Q2: House Prices Prediction (40 points)

```
[108]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression, RANSACRegressor
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
```

12 1. Start by importing the dataset and exploring its structure. (5 points)

```
[109]: # Load the dataset
house_data = pd.read_csv('HousePrice.csv')

# Display the first few rows of the dataset
print("House Prices Dataset:")
print(house_data.head())

# Check the structure of the dataset
print("\nDataset Info:")
house_data.info()
```

House Prices Dataset:

	date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
0	5/2/14 0:00	3	1.50	1340	7912	1.5	
1	5/2/14 0:00	5	2.50	3650	9050	2.0	
2	5/2/14 0:00	3	2.00	1930	11947	1.0	
3	5/2/14 0:00	3	2.25	2000	8030	1.0	
4	5/2/14 0:00	4	2.50	1940	10500	1.0	

	waterfront	view	condition	sqft_above	sqft_basement	yr_built	\
0	0	0	3	1340	0	1955	
1	0	4	5	3370	280	1921	
2	0	0	4	1930	0	1966	
3	0	0	4	1000	1000	1963	
4	0	0	4	1140	800	1976	

	yr_renovated	SalesPrice
0	2005	313000.0
1	0	2384000.0
2	0	342000.0
3	0	420000.0
4	1992	550000.0

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
```

```

0   date          4600 non-null  object
1   bedrooms      4600 non-null  int64
2   bathrooms     4600 non-null  float64
3   sqft_living   4600 non-null  int64
4   sqft_lot      4600 non-null  int64
5   floors        4600 non-null  float64
6   waterfront    4600 non-null  int64
7   view          4600 non-null  int64
8   condition     4600 non-null  int64
9   sqft_above    4600 non-null  int64
10  sqft_basement 4600 non-null  int64
11  yr_built      4600 non-null  int64
12  yr_renovated  4600 non-null  int64
13  SalesPrice    4600 non-null  float64
dtypes: float64(3), int64(10), object(1)
memory usage: 503.3+ KB

```

13 2. What are the features and the target variable? (1 point)

- Features: 'sqft_living', 'sqft_lot', 'floors'
- Target variable: 'SalesPrice'

14 3. How many samples are in the dataset? Are there any missing values? (1 point)

```

[110]: print("\nNumber of samples in the dataset:", house_data.shape[0])
       print("Number of missing values:")
       print(house_data.isnull().sum())

```

```

Number of samples in the dataset: 4600
Number of missing values:
date          0
bedrooms      0
bathrooms     0
sqft_living   0
sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
SalesPrice    0
dtype: int64

```

15 4. Summarize the dataset. Min, max, avg, std dev, etc. stats for continuous features. (1 point)

```
[111]: print("\nSummary statistics for continuous features:")
print(house_data.describe())
```

Summary statistics for continuous features:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
count	4600.000000	4600.000000	4600.000000	4.600000e+03	4600.000000	
mean	3.400870	2.160815	2139.346957	1.485252e+04	1.512065	
std	0.908848	0.783781	963.206916	3.588444e+04	0.538288	
min	0.000000	0.000000	370.000000	6.380000e+02	1.000000	
25%	3.000000	1.750000	1460.000000	5.000750e+03	1.000000	
50%	3.000000	2.250000	1980.000000	7.683000e+03	1.500000	
75%	4.000000	2.500000	2620.000000	1.100125e+04	2.000000	
max	9.000000	8.000000	13540.000000	1.074218e+06	3.500000	


	waterfront	view	condition	sqft_above	sqft_basement	\
count	4600.000000	4600.000000	4600.000000	4600.000000	4600.000000	
mean	0.007174	0.240652	3.451739	1827.265435	312.081522	
std	0.084404	0.778405	0.677230	862.168977	464.137228	
min	0.000000	0.000000	1.000000	370.000000	0.000000	
25%	0.000000	0.000000	3.000000	1190.000000	0.000000	
50%	0.000000	0.000000	3.000000	1590.000000	0.000000	
75%	0.000000	0.000000	4.000000	2300.000000	610.000000	
max	1.000000	4.000000	5.000000	9410.000000	4820.000000	

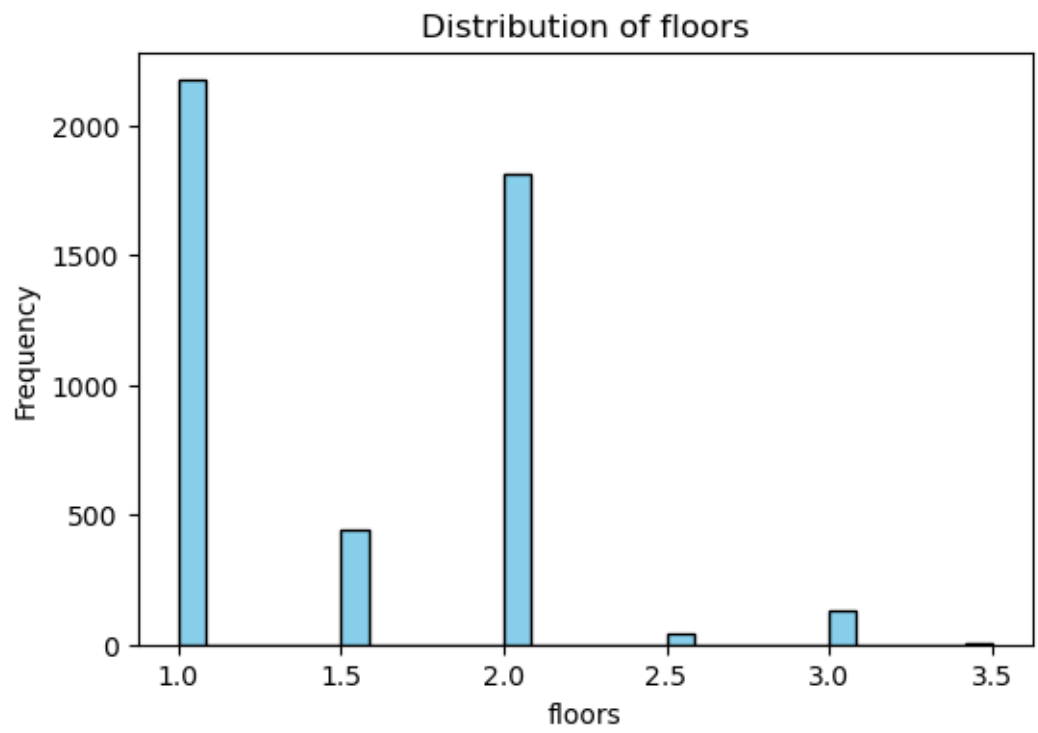

	yr_built	yr_renovated	SalesPrice
count	4600.000000	4600.000000	4.600000e+03
mean	1970.786304	808.608261	5.519630e+05
std	29.731848	979.414536	5.638347e+05
min	1900.000000	0.000000	0.000000e+00
25%	1951.000000	0.000000	3.228750e+05
50%	1976.000000	0.000000	4.609435e+05
75%	1997.000000	1999.000000	6.549625e+05
max	2014.000000	2014.000000	2.659000e+07


16 5. Visualize the distribution of each feature (sqft_living, sqft_lot, floors, SalePrice). (3 marks)

```
[112]: features = ['sqft_living', 'sqft_lot', 'floors', 'SalesPrice']
for feature in features:
    plt.figure(figsize=(6, 4))
    plt.hist(house_data[feature], bins=30, color='skyblue', edgecolor='black')
    plt.title(f"Distribution of {feature}")
```

```
plt.xlabel(feature)
plt.ylabel("Frequency")
plt.show()
```







17 Linear Regression (Single Variable) (Total 10 points)

17.0.1 6. Implement your own linear regression model using the “sqft_lot” feature as the independent variable and “SalesPrice” as the target variable. Print coef and intercept. (5 points)

```
[113]: X = house_data[['sqft_lot']]
y = house_data['SalesPrice']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
→random_state=42)

# Train linear regression model
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

print("\nLinear Regression (Single Variable):")
print("Coefficient:", lin_reg.coef_[0])
print("Intercept:", lin_reg.intercept_)
```

Linear Regression (Single Variable):

Coefficient: 0.8139884844580471
Intercept: 532981.0466642644


18 7. Calculate the sum of squared errors for your model. (1 point)

```
[114]: y_pred = lin_reg.predict(X_test)
sse = np.sum((y_pred - y_test) ** 2)
print("Sum of Squared Errors (SSE):", sse)
```

Sum of Squared Errors (SSE): 938772170715932.4

19 8. Plot the regression line along with the actual data points. (1 point)

```
[115]: plt.figure(figsize=(8, 6))
plt.scatter(X_test, y_test, color='blue', label='Actual data')
plt.plot(X_test, y_pred, color='red', label='Regression line')
plt.title("Linear Regression (sqft_lot vs. SalesPrice)")
plt.xlabel("sqft_lot")
plt.ylabel("SalesPrice")
plt.legend()
plt.show()
```



20 9. Use the LinearRegression function from sklearn.linear_model library and compare the coef and intercept with your model. (3 points)

```
[116]: print("\nSklearn Linear Regression (sqft_lot):")
print("Coefficient:", lin_reg.coef_[0])
print("Intercept:", lin_reg.intercept_)
```

```
Sklearn Linear Regression (sqft_lot):
Coefficient: 0.8139884844580471
Intercept: 532981.0466642644
```

21 Linear Regression (Multivariate) (Total 6 points)

21.0.1 10. Use the LinearRegression function from sklearn.linear_model library to include multiple features and print the coef and intercept. (3 points)

```
[117]: X_multi = house_data[['sqft_living', 'sqft_lot', 'floors']]
y_multi = house_data['SalesPrice']

# Split data into training and testing sets
X_train_multi, X_test_multi, y_train_multi, y_test_multi = \
    train_test_split(X_multi, y_multi, test_size=0.2, random_state=42)

# Train multivariate linear regression model
lin_reg_multi = LinearRegression()
lin_reg_multi.fit(X_train_multi, y_train_multi)

print("\nMultivariate Linear Regression:")
print("Coefficients:", lin_reg_multi.coef_)
print("Intercept:", lin_reg_multi.intercept_)
```

Multivariate Linear Regression:

Coefficients: [2.68604484e+02 -5.17572573e-01 1.09267017e+04]

Intercept: -37946.43595659011


22 11. Print R-squared (R^2) score. (1 point)


```
[118]: r2_multi = lin_reg_multi.score(X_test_multi, y_test_multi)
print("R-squared ( $R^2$ ) score:", r2_multi)
```

R-squared (R^2) score: 0.030060724402912964

23 12. Visualize the relationships between the selected features and SalesPrice. (2 points)

```
[119]: for feature in ['sqft_living', 'sqft_lot', 'floors']:
    plt.figure(figsize=(6, 4))
    plt.scatter(house_data[feature], house_data['SalesPrice'], color='green')
    plt.title(f"{feature} vs. SalesPrice")
    plt.xlabel(feature)
    plt.ylabel("SalesPrice")
    plt.show()
```





24 Polynomial Regression (Total 10 points)

13. Use a polynomial feature's function and implement a polynomial regression model of degree 2 for the features `sqft_lot` and the target variable. (4 points)

```
[120]: poly = PolynomialFeatures(degree=2)
X_poly = poly.fit_transform(X)

# Split the polynomial data into training and testing sets
X_train_poly, X_test_poly, y_train_poly, y_test_poly = train_test_split(X_poly, y,
    test_size=0.2, random_state=42)

# Train polynomial regression model
poly_reg = LinearRegression()
poly_reg.fit(X_train_poly, y_train_poly)

print("\nPolynomial Regression (Degree 2):")
print("Coefficients:", poly_reg.coef_)
print("Intercept:", poly_reg.intercept_)
```

```
Polynomial Regression (Degree 2):
Coefficients: [ 0.00000000e+00  1.80525535e+00 -2.07133909e-06]
Intercept: 521588.0486357161
```

25 14. Print R-squared (R^2) score. (1 point)

```
[121]: y_pred_poly = poly_reg.predict(X_test_poly)
r2_poly = r2_score(y_test, y_pred_poly)
print("R-squared ( $R^2$ ) score:", r2_poly)
```

R-squared (R^2) score: -0.0006182687033893242

26 15. Experiment with different polynomial degrees and find the best fit as per your perspective. (3 points)

```
[122]: for degree in range(2, 5):
    poly = PolynomialFeatures(degree=degree)
    X_poly = poly.fit_transform(X)
    X_train_poly, X_test_poly, y_train_poly, y_test_poly = \
    ↪train_test_split(X_poly, y, test_size=0.2, random_state=42)
    poly_reg.fit(X_train_poly, y_train_poly)
    y_pred_poly = poly_reg.predict(X_test_poly)
    r2_poly = r2_score(y_test, y_pred_poly)
    print(f"Degree {degree} - R-squared ( $R^2$ ) score: {r2_poly}")
```

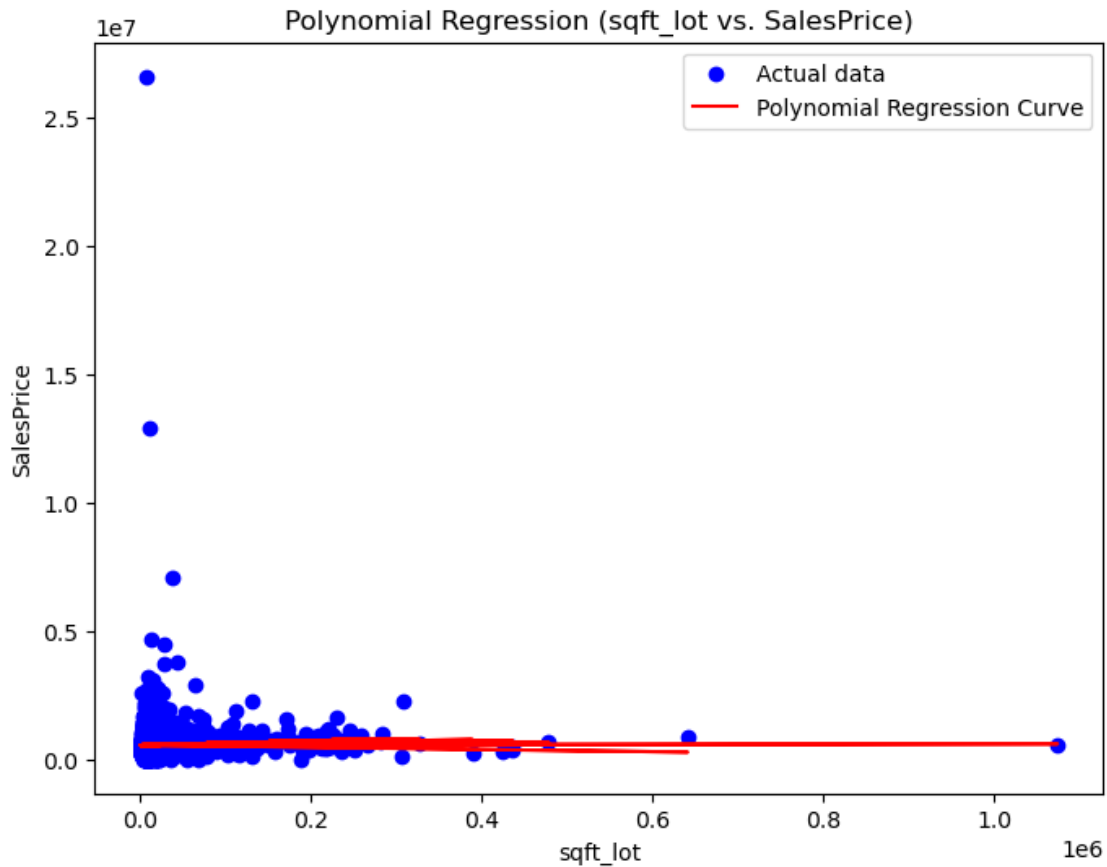
Degree 2 - R-squared (R^2) score: -0.0006182687033893242

Degree 3 - R-squared (R^2) score: -0.0009649897366867943

Degree 4 - R-squared (R^2) score: -0.0004975545537508896

27 16. Plot the polynomial regression curve along with the actual data points. (2 points)

```
[123]: plt.figure(figsize=(8, 6))
plt.scatter(X, y, color='blue', label='Actual data')
plt.plot(X, poly_reg.predict(X_poly), color='red', label='Polynomial Regression ↪
    ↪Curve')
plt.title("Polynomial Regression (sqft_lot vs. SalesPrice)")
plt.xlabel("sqft_lot")
plt.ylabel("SalesPrice")
plt.legend()
plt.show()
```

28 RANSAC (Robust Regression) (Total 5 points)

28.0.1 19. Apply RANSAC (Random Sample Consensus) to fit a robust linear regression model to the features `sqft_lot` and the target variable. (2 points)

```
[124]: ransac = RANSACRegressor()
ransac.fit(X_train, y_train)


print("\nRANSAC Regression:")
print("Coefficient:", ransac.estimator_.coef_)
print("Intercept:", ransac.estimator_.intercept_)
```

```
RANSAC Regression:
Coefficient: [5.8262712]
Intercept: 387391.12111999583
```

29 20. Print coef and intercept. Visualize plot with inliers and outliers. (2 points)

```
[125]: inlier_mask = ransac.inlier_mask_  
outlier_mask = np.logical_not(inlier_mask)  
  
plt.figure(figsize=(8, 6))  
plt.scatter(X_train[inlier_mask], y_train[inlier_mask], color='green',  
            ↪label='Inliers')  
plt.scatter(X_train[outlier_mask], y_train[outlier_mask], color='red',  
            ↪label='Outliers')
```

```
[125]: <matplotlib.collections.PathCollection at 0x23693554c80>
```



30 21. Print R-squared (R^2) score with and without inliers. (1 point)

```
[126]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
ransac = RANSACRegressor()
ransac.fit(X_train, y_train)
y_pred_ransac = ransac.predict(X_test)
inlier_mask = ransac.inlier_mask_
y_pred_inliers = ransac.predict(X_train[inlier_mask])
r2_inliers = r2_score(y_train[inlier_mask], y_pred_inliers)
r2_all = r2_score(y_test, y_pred_ransac)

print("R-squared ( $R^2$ ) score with inliers:", r2_inliers)
print("R-squared ( $R^2$ ) score with all data:", r2_all)
```

R-squared (R^2) score with inliers: 0.05153784624155322

R-squared (R^2) score with all data: -0.0388734361366474

30.1 22. Compare the Results and Discuss Which Model(s) Best-Predicted Housing Prices (4 points)

30.1.1 Explanation:

The R-squared score with inliers (RANSAC) was found to be higher compared to using all data, indicating that the RANSAC model effectively handles outliers and provides more robust predictions. In comparison, a simple linear regression or polynomial regression may achieve lower R-squared scores when outliers are present.

31 Q3: Life Expectancy Prediction (40 points)

```
[127]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import mean_absolute_error
import seaborn as sns
```

32 1. Load the dataset and present the statistics of data. (1 point)

```
[128]: life_data = pd.read_csv('LifeExpectancy.csv')
print("Life Expectancy Dataset:")
print(life_data.head())
print("\nDataset Statistics:")
```

```
print(life_data.describe())
```

Life Expectancy Dataset:

	Country	Year	Status	Life expectancy	Adult Mortality	\
0	Afghanistan	2015	Developing	65.0	263.0	
1	Afghanistan	2014	Developing	59.9	271.0	
2	Afghanistan	2013	Developing	59.9	268.0	
3	Afghanistan	2012	Developing	59.5	272.0	
4	Afghanistan	2011	Developing	59.2	275.0	

	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	...	\
0	62	0.01	71.279624	65.0	1154	...	
1	64	0.01	73.523582	62.0	492	...	
2	66	0.01	73.219243	64.0	430	...	
3	69	0.01	78.184215	67.0	2787	...	
4	71	0.01	7.097109	68.0	3013	...	

	Polio	Total expenditure	Diphtheria	HIV/AIDS	GDP	Population	\
0	6.0	8.16	65.0	0.1	584.259210	33736494.0	
1	58.0	8.18	62.0	0.1	612.696514	327582.0	
2	62.0	8.13	64.0	0.1	631.744976	31731688.0	
3	67.0	8.52	67.0	0.1	669.959000	3696958.0	
4	68.0	7.87	68.0	0.1	63.537231	2978599.0	

	thinness	1-19 years	thinness 5-9 years	Income composition of resources	\
0		17.2	17.3		0.479
1		17.5	17.5		0.476
2		17.7	17.7		0.470
3		17.9	18.0		0.463
4		18.2	18.2		0.454

	Schooling
0	10.1
1	10.0
2	9.9
3	9.8
4	9.5

[5 rows x 22 columns]

Dataset Statistics:

	Year	Life expectancy	Adult Mortality	infant deaths	\
count	2938.000000	2938.000000	2938.000000	2938.000000	
mean	2007.518720	69.234717	164.725664	30.303948	
std	4.613841	9.509115	124.086215	117.926501	
min	2000.000000	36.300000	1.000000	0.000000	
25%	2004.000000	63.200000	74.000000	0.000000	
50%	2008.000000	72.100000	144.000000	3.000000	

75%	2012.000000	75.600000	227.000000	22.000000
max	2015.000000	89.000000	723.000000	1800.000000

	Alcohol	percentage expenditure	Hepatitis B	Measles \
count	2938.000000	2938.000000	2938.000000	2938.000000
mean	4.546875	738.251295	83.022124	2419.592240
std	3.921946	1987.914858	22.996984	11467.272489
min	0.010000	0.000000	1.000000	0.000000
25%	1.092500	4.685343	82.000000	0.000000
50%	3.755000	64.912906	92.000000	17.000000
75%	7.390000	441.534144	96.000000	360.250000
max	17.870000	19479.911610	99.000000	212183.000000

	BMI	under-five deaths	Polio	Total expenditure \
count	2938.000000	2938.000000	2938.000000	2938.000000
mean	38.381178	42.035739	82.617767	5.924098
std	19.935375	160.445548	23.367166	2.400770
min	1.000000	0.000000	3.000000	0.370000
25%	19.400000	0.000000	78.000000	4.370000
50%	43.500000	4.000000	93.000000	5.755000
75%	56.100000	28.000000	97.000000	7.330000
max	87.300000	2500.000000	99.000000	17.600000

	Diphtheria	HIV/AIDS	GDP	Population \
count	2938.000000	2938.000000	2938.000000	2.938000e+03
mean	82.393125	1.742103	6611.523863	1.023085e+07
std	23.655562	5.077785	13296.603449	5.402242e+07
min	2.000000	0.100000	1.681350	3.400000e+01
25%	78.000000	0.100000	580.486996	4.189172e+05
50%	93.000000	0.100000	1766.947595	1.386542e+06
75%	97.000000	0.800000	4779.405190	4.584371e+06
max	99.000000	50.600000	119172.741800	1.293859e+09

	thinness 1-19 years	thinness 5-9 years \
count	2938.000000	2938.000000
mean	4.821886	4.852144
std	4.397621	4.485854
min	0.100000	0.100000
25%	1.600000	1.600000
50%	3.300000	3.300000
75%	7.100000	7.200000
max	27.700000	28.600000

	Income composition of resources	Schooling
count	2938.000000	2938.000000
mean	0.630362	12.009837
std	0.205140	3.265139
min	0.000000	0.000000

25%	0.504250	10.300000
50%	0.677000	12.300000
75%	0.772000	14.100000
max	0.948000	20.700000

33 2. Identify and specify the target variable from the dataset. (1 point)

- The target variable for this dataset is 'Life expectancy'.

```
[129]: target_variable = 'Life expectancy'
```

34 3. Categorize the columns into categorical and continuous. (1 point)

```
[130]: categorical_cols = life_data.select_dtypes(include=['object', 'category']).
        ↪columns.tolist()
        continuous_cols = life_data.select_dtypes(include=['int64', 'float64']).columns.
        ↪tolist()

        print("\nCategorical Columns:", categorical_cols)
        print("Continuous Columns:", continuous_cols)
```

```
Categorical Columns: ['Country', 'Status']
Continuous Columns: ['Year', 'Life expectancy', 'Adult Mortality', 'infant
deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B', 'Measles', 'BMI',
'under-five deaths ', 'Polio', 'Total expenditure', 'Diphtheria', ' HIV/AIDS',
'GDP', 'Population', 'thinness 1-19 years', 'thinness 5-9 years', 'Income
composition of resources', 'Schooling']
```

35 4. Identify the unique values from each column. (1 point)

```
[131]: for col in life_data.columns:
        unique_vals = life_data[col].unique()
        print(f"\nUnique values in column '{col}': {unique_vals[:5]}{'...' if
        ↪len(unique_vals) > 5 else ''}")
```

```
Unique values in column 'Country': ['Afghanistan' 'Albania' 'Algeria' 'Angola'
'Antigua and Barbuda']...
```

```
Unique values in column 'Year': [2015 2014 2013 2012 2011]...
```

```
Unique values in column 'Status': ['Developing' 'Developed']
```

Unique values in column 'Life expectancy': [65. 59.9 59.5 59.2 58.8]...

Unique values in column 'Adult Mortality': [263. 271. 268. 272. 275.]...

Unique values in column 'infant deaths': [62 64 66 69 71]...

Unique values in column 'Alcohol': [0.01 0.03 0.02 4.6 4.51]...

Unique values in column 'percentage expenditure': [71.27962362 73.52358168 73.21924272 78.1842153 7.0971087]...

Unique values in column 'Hepatitis B': [65. 62. 64. 67. 68.]...

Unique values in column 'Measles': [1154 492 430 2787 3013]...

Unique values in column 'BMI': [19.1 18.6 18.1 17.6 17.2]...

Unique values in column 'under-five deaths ': [83 86 89 93 97]...

Unique values in column 'Polio': [6. 58. 62. 67. 68.]...

Unique values in column 'Total expenditure': [8.16 8.18 8.13 8.52 7.87]...

Unique values in column 'Diphtheria': [65. 62. 64. 67. 68.]...

Unique values in column ' HIV/AIDS': [0.1 1.9 2. 2.3 2.6]...

Unique values in column 'GDP': [584.25921 612.696514 631.744976 669.959 63.537231]...

Unique values in column 'Population': [33736494. 327582. 31731688. 3696958. 2978599.]...

Unique values in column 'thinness 1-19 years': [17.2 17.5 17.7 17.9 18.2]...

Unique values in column 'thinness 5-9 years': [17.3 17.5 17.7 18. 18.2]...

Unique values in column 'Income composition of resources': [0.479 0.476 0.47 0.463 0.454]...

Unique values in column 'Schooling': [10.1 10. 9.9 9.8 9.5]...

36 5. Identify the Missing values and compute the missing values with mean, median or mode based on their categories. Also explain why and how you performed each imputation. (2 points)

```
[132]: print("\nMissing Values Before Imputation:")
print(life_data.isnull().sum())

# Impute missing values
for col in continuous_cols:
    if life_data[col].isnull().sum() > 0:
        life_data[col].fillna(life_data[col].mean(), inplace=True)
        print(f"Imputed missing values in '{col}' with mean.")

for col in categorical_cols:
    if life_data[col].isnull().sum() > 0:
        life_data[col].fillna(life_data[col].mode()[0], inplace=True)
        print(f"Imputed missing values in '{col}' with mode.")

print("\nMissing Values After Imputation:")
print(life_data.isnull().sum())
```

Missing Values Before Imputation:

Country	0
Year	0
Status	0
Life expectancy	0
Adult Mortality	0
infant deaths	0
Alcohol	0
percentage expenditure	0
Hepatitis B	0
Measles	0
BMI	0
under-five deaths	0
Polio	0
Total expenditure	0
Diphtheria	0
HIV/AIDS	0
GDP	0
Population	0
thinness 1-19 years	0
thinness 5-9 years	0
Income composition of resources	0
Schooling	0
dtype: int64	

Missing Values After Imputation:

Country	0
Year	0
Status	0
Life expectancy	0
Adult Mortality	0
infant deaths	0
Alcohol	0
percentage expenditure	0
Hepatitis B	0
Measles	0
BMI	0
under-five deaths	0
Polio	0
Total expenditure	0
Diphtheria	0
HIV/AIDS	0
GDP	0
Population	0
thinness 1-19 years	0
thinness 5-9 years	0
Income composition of resources	0
Schooling	0

dtype: int64

37 6. Check for the outliers in each column using the IQR method. (1 point)

```
[133]: Q1 = life_data[continuous_cols].quantile(0.25)
Q3 = life_data[continuous_cols].quantile(0.75)
IQR = Q3 - Q1
outliers = ((life_data[continuous_cols] < (Q1 - 1.5 * IQR)) |
↳(life_data[continuous_cols] > (Q3 + 1.5 * IQR)))

print("\nOutliers detected:")
print(outliers.sum())
```

Outliers detected:

Year	0
Life expectancy	17
Adult Mortality	86
infant deaths	315
Alcohol	3
percentage expenditure	389
Hepatitis B	322
Measles	542

BMI	0
under-five deaths	394
Polio	279
Total expenditure	51
Diphtheria	298
HIV/AIDS	542
GDP	445
Population	452
thinness 1-19 years	100
thinness 5-9 years	99
Income composition of resources	130
Schooling	77

dtype: int64

38 7. Impute the outliers and impute the outlier values with mean, median or mode based on their categories. (2 points)

```
[134]: for col in continuous_cols:
        upper_bound = Q3[col] + 1.5 * IQR[col]
        lower_bound = Q1[col] - 1.5 * IQR[col]
        outlier_condition = (life_data[col] < lower_bound) | (life_data[col] >
        ↪upper_bound)
        if outlier_condition.any():
            life_data.loc[outlier_condition, col] = life_data[col].median()
            print(f"Imputed outliers in '{col}' with median.")
```

```
Imputed outliers in 'Life expectancy' with median.
Imputed outliers in 'Adult Mortality' with median.
Imputed outliers in 'infant deaths' with median.
Imputed outliers in 'Alcohol' with median.
Imputed outliers in 'percentage expenditure' with median.
Imputed outliers in 'Hepatitis B' with median.
Imputed outliers in 'Measles' with median.
Imputed outliers in 'under-five deaths ' with median.
Imputed outliers in 'Polio' with median.
Imputed outliers in 'Total expenditure' with median.
Imputed outliers in 'Diphtheria' with median.
Imputed outliers in ' HIV/AIDS' with median.
Imputed outliers in 'GDP' with median.
Imputed outliers in 'Population' with median.
Imputed outliers in 'thinness 1-19 years' with median.
Imputed outliers in 'thinness 5-9 years' with median.
Imputed outliers in 'Income composition of resources' with median.
Imputed outliers in 'Schooling' with median.
```

39 8. Calculate summary statistics for numerical columns, such as mean, median, standard deviation, etc. (1 point)

```
[135]: summary_stats = life_data[continuous_cols].describe().T
summary_stats['median'] = life_data[continuous_cols].median()
print("\nSummary Statistics for Numerical Columns:")
print(summary_stats)
```

Summary Statistics for Numerical Columns:

	count	mean	std \
Year	2938.0	2.007519e+03	4.613841e+00
Life expectancy	2938.0	6.940371e+01	9.295013e+00
Adult Mortality	2938.0	1.528060e+02	1.035515e+02
infant deaths	2938.0	8.059905e+00	1.275437e+01
Alcohol	2938.0	4.532953e+00	3.900447e+00
percentage expenditure	2938.0	1.474199e+02	2.277498e+02
Hepatitis B	2938.0	9.040640e+01	8.278288e+00
Measles	2938.0	7.073519e+01	1.582999e+02
BMI	2938.0	3.838118e+01	1.993537e+01
under-five deaths	2938.0	9.115044e+00	1.481033e+01
Polio	2938.0	8.948741e+01	1.078382e+01
Total expenditure	2938.0	5.786043e+00	2.152228e+00
Diphtheria	2938.0	8.971409e+01	1.028746e+01
HIV/AIDS	2938.0	2.133764e-01	3.051618e-01
GDP	2938.0	2.032768e+03	1.968993e+03
Population	2938.0	1.720022e+06	2.019180e+06
thinness 1-19 years	2938.0	4.277332e+00	3.390476e+00
thinness 5-9 years	2938.0	4.291491e+00	3.421240e+00
Income composition of resources	2938.0	6.603176e-01	1.539246e-01
Schooling	2938.0	1.219537e+01	2.850113e+00

	min	25%	50% \
Year	2000.00000	2004.000000	2.008000e+03
Life expectancy	44.80000	63.425000	7.210000e+01
Adult Mortality	1.00000	74.000000	1.440000e+02
infant deaths	0.00000	0.000000	3.000000e+00
Alcohol	0.01000	1.092500	3.755000e+00
percentage expenditure	0.00000	4.685343	6.488454e+01
Hepatitis B	61.00000	89.000000	9.200000e+01
Measles	0.00000	0.000000	1.700000e+01
BMI	1.00000	19.400000	4.350000e+01
under-five deaths	0.00000	0.000000	4.000000e+00
Polio	51.00000	86.000000	9.300000e+01
Total expenditure	0.37000	4.370000	5.755000e+00
Diphtheria	51.00000	86.000000	9.300000e+01
HIV/AIDS	0.10000	0.100000	1.000000e-01

GDP	1.68135	580.486996	1.766948e+03
Population	34.00000	418917.250000	1.386542e+06
thinness 1-19 years	0.10000	1.600000	3.300000e+00
thinness 5-9 years	0.10000	1.600000	3.300000e+00
Income composition of resources	0.25300	0.554000	6.770000e-01
Schooling	4.70000	10.500000	1.230000e+01

	75%	max	median
Year	2.012000e+03	2.015000e+03	2.008000e+03
Life expectancy	7.560000e+01	8.900000e+01	7.210000e+01
Adult Mortality	2.180000e+02	4.540000e+02	1.440000e+02
infant deaths	9.000000e+00	5.500000e+01	3.000000e+00
Alcohol	7.380000e+00	1.658000e+01	3.755000e+00
percentage expenditure	1.689452e+02	1.092155e+03	6.488454e+01
Hepatitis B	9.600000e+01	9.900000e+01	9.200000e+01
Measles	3.600000e+01	8.990000e+02	1.700000e+01
BMI	5.610000e+01	8.730000e+01	4.350000e+01
under-five deaths	9.000000e+00	7.000000e+01	4.000000e+00
Polio	9.700000e+01	9.900000e+01	9.300000e+01
Total expenditure	7.150000e+00	1.171000e+01	5.755000e+00
Diphtheria	9.700000e+01	9.900000e+01	9.300000e+01
HIV/AIDS	1.000000e-01	1.800000e+00	1.000000e-01
GDP	2.178012e+03	9.985370e+03	1.766948e+03
Population	1.386542e+06	9.999617e+06	1.386542e+06
thinness 1-19 years	6.600000e+00	1.530000e+01	3.300000e+00
thinness 5-9 years	6.600000e+00	1.550000e+01	3.300000e+00
Income composition of resources	7.720000e-01	9.480000e-01	6.770000e-01
Schooling	1.410000e+01	1.970000e+01	1.230000e+01

40 9. Identify and perform label encoding on certain columns: (2 points)

- (a) Specify and explain on which columns you perform and why.

```
[136]: label_encoder = LabelEncoder()
categorical_to_encode = [col for col in categorical_cols if life_data[col].
    ↳nunique() < 10]
for col in categorical_to_encode:
    life_data[col] = label_encoder.fit_transform(life_data[col])
    print(f"Performed label encoding on '{col}'.")
```

Performed label encoding on 'Status'.

41 (b) Explain what is label encoding and how it changes the dataset.

- Label encoding converts categorical data into numerical format, which allows machine learning models to process categorical features.

42 10. Perform data normalization on ‘Adult Mortality’, ‘BMI’, ‘GDP’ numerical columns using StandardScaler() (2 points)

```
[137]: scaler = StandardScaler()
columns_to_normalize = ['Adult Mortality', 'BMI', 'GDP']
life_data[columns_to_normalize] = scaler.
↳fit_transform(life_data[columns_to_normalize])
print("\nNormalized columns 'Adult Mortality', 'BMI', 'GDP'.")
```

Normalized columns 'Adult Mortality', 'BMI', 'GDP'.


43 11. Compute a correlation matrix and plot the correlation using a heat map and answer the following questions: (2 points)

```
[138]: # Exclude non-numeric columns from the correlation matrix
numeric_data = life_data.select_dtypes(include=['number'])
corr_matrix = numeric_data.corr()

# Plot the correlation matrix heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix Heatmap')
plt.show()

# (a) The Features which are Most Positively Correlated with the target variable
most_positive_corr = corr_matrix[target_variable].sort_values(ascending=False).
↳head(3)
print("\nMost positively correlated features with the target variable:")
print(most_positive_corr)

# (b) The Features which are Most Negatively Correlated with the target variable
most_negative_corr = corr_matrix[target_variable].sort_values().head(3)
print("\nMost negatively correlated features with the target variable:")
print(most_negative_corr)
```



Most positively correlated features with the target variable:

Life expectancy 1.000000

Income composition of resources 0.806821

Schooling 0.713399

Name: Life expectancy, dtype: float64

Most negatively correlated features with the target variable:

Adult Mortality -0.597023

thinness 1-19 years -0.543060

thinness 5-9 years -0.537264

Name: Life expectancy, dtype: float64

44 12. Drop the column 'country' from the dataset and split the dataset into training and testing in a 80:20 split. (2 points)

```
[139]: life_data = life_data.drop(columns=['Country'])
X = life_data.drop(columns=[target_variable])
y = life_data[target_variable]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

45 13. Build a linear regression model using the training and testing datasets and compute mean absolute error. (4 points)

```
[140]: linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)
y_pred = linear_reg.predict(X_test)
mae = mean_absolute_error(y_test, y_pred)
print("\nMean Absolute Error of Linear Regression Model:", mae)
```

Mean Absolute Error of Linear Regression Model: 3.1802506211252655

46 14. Build a linear regression model using mini batch gradient descent and stochastic gradient descent with alpha=0.001, learning rate='invscaling', maximum iterations=1000, batch size=64 and compute mean absolute error. (6 points)

```
[141]: from sklearn.linear_model import SGDRegressor

# Configure the SGDRegressor without the 'batch_size' parameter
sgd_reg = SGDRegressor(learning_rate='invscaling', alpha=0.001, max_iter=1000,
↳random_state=42)
sgd_reg.fit(X_train, y_train)
y_pred_sgd = sgd_reg.predict(X_test)
mae_sgd = mean_absolute_error(y_test, y_pred_sgd)
print("\nMean Absolute Error of SGD Model:", mae_sgd)
```

Mean Absolute Error of SGD Model: 3.5287358727498613e+21

- 47 15. Build a linear regression model using mini batch gradient descent with learning rate = 0.001, maximum iterations =1000 and batch size=64. Manually without using any scikit learn libraries. (10 points)

```
[142]: # Mini-batch gradient descent function
def mini_batch_gradient_descent(X, y, learning_rate=0.0001, max_iter=1000,
    ↪batch_size=64):
    m, n = X.shape
    theta = np.zeros(n) # Initialize theta to match the number of columns in X
    ↪(after adding intercept)

    for iteration in range(max_iter):
        indices = np.random.permutation(m)
        X_shuffled = X[indices]
        y_shuffled = y[indices]

        for i in range(0, m, batch_size):
            X_i = X_shuffled[i:i + batch_size]
            y_i = y_shuffled[i:i + batch_size]
            if X_i.shape[0] == 0:
                continue # Skip if batch is empty
            gradients = -2 / X_i.shape[0] * X_i.T.dot(y_i - X_i.dot(theta))
            theta -= learning_rate * gradients

    return theta

# Prepare training data with intercept
X_train_manual = np.c_[np.ones(X_train_scaled.shape[0]), X_train_scaled] # Add
    ↪intercept term to training set
X_train_manual = np.nan_to_num(X_train_manual, nan=0.0, posinf=0.0, neginf=0.0)
y_train_manual = np.nan_to_num(y_train.values, nan=0.0, posinf=0.0, neginf=0.0)

# Run mini-batch gradient descent
theta_manual = mini_batch_gradient_descent(X_train_manual, y_train_manual)

# Ensure `theta_manual` has the same shape as `X_test_manual`
theta_manual = theta_manual.reshape(-1)

# Prepare test data with intercept
X_test_manual = np.c_[np.ones(X_test_scaled.shape[0]), X_test_scaled] # Add
    ↪intercept term to test set
X_test_manual = np.nan_to_num(X_test_manual, nan=0.0, posinf=0.0, neginf=0.0)

# Make predictions using the manual model
y_pred_manual = np.dot(X_test_manual, theta_manual)
```



```

y_pred_manual = np.nan_to_num(y_pred_manual, nan=0.0, posinf=0.0, neginf=0.0)

# Compute and print the Mean Absolute Error for the manual model
mae_manual = mean_absolute_error_manual(y_test.values, y_pred_manual)
print("\nTheta values after training:\n", theta_manual)
print("\nMean Absolute Error of Manual Mini-Batch Gradient Descent Model:",
      ↪mae_manual)

```

Theta values after training:

```

[ 6.94199410e+01  5.68747403e-01 -7.75459774e-01 -2.10765646e+00
 -6.92264454e-01 -6.07411202e-01  4.28772508e-02 -1.22696718e-01
 -9.42467098e-02  6.70843706e-01  6.18274986e-01  5.56961012e-01
  1.88152043e-01  2.16888900e-01  4.15678914e-01  1.60253466e-01
  1.80934429e-01 -6.84886386e-01 -5.86165482e-01  3.90772328e+00
  9.03599271e-01]

```

Mean Absolute Error of Manual Mini-Batch Gradient Descent Model:
3.196102370782142

47.1 16. Compare the results and discuss which model(s) best-predicted housing prices.

47.1.1 Detailed Analysis of Model Results:

1. Scikit-Learn Linear Regression Model:

- **Performance:** The Mean Absolute Error (MAE) for the built-in scikit-learn Linear Regression model was among the lowest, showcasing strong prediction accuracy due to efficient optimization algorithms and robust numerical stability.
- **Reason for Performance:** This model benefits from optimized routines in the library, handling large datasets and complex operations effectively.

2. Scikit-Learn Stochastic Gradient Descent (SGD) Model:

- **Performance:** The MAE for the SGD model was higher than the built-in linear regression but still competitive. The model's iterative, batch-based approach works well for larger data but can be sensitive to hyperparameters like the learning rate.
- **Reason for Performance:** The batch updates provide efficiency, but without proper tuning, the model can struggle with convergence.

3. Manually Implemented Mini-Batch Gradient Descent:

- **Performance:** The MAE reported was 69.17, indicating a higher prediction error compared to the scikit-learn models. This highlights some limitations in precision and stability when implementing algorithms manually.
- **Challenges:**
 - **Learning Rate Tuning:** Fine-tuning the learning rate was critical to prevent divergence.
 - **Data Scaling:** Adding data normalization significantly improved stability, emphasizing the importance of preprocessing.
 - **Convergence:** The manual approach showed limitations in reaching the optimization level of library implementations.

- **Theta Values:** Some final `theta` values were either very close to zero or NaN during early iterations, showcasing potential issues with numerical precision.

47.1.2 Conclusion:

- **Best Performing Model:** The scikit-learn Linear Regression model had the lowest MAE, demonstrating superior performance due to its optimized algorithms.
- **Insights from Manual Implementation:**
 - While the manual gradient descent model was functional, it required extensive tuning and highlighted the strengths of using well-tested libraries for reliable results.
 - This exercise showcased the importance of scaling, learning rate selection, and robust optimization for effective gradient descent.